

Sukhrob Ilyosbekov

Senior Software Engineer · AI/ML & Computer Vision

📍 Boston, MA 📞 (857) 867-1942 ✉ silyosbekov@gmail.com 🌐 suxrobgm.net
🌐 [linkedin.com/in/suxrobgm](https://www.linkedin.com/in/suxrobgm) 🐙 github.com/suxrobgm

SUMMARY

Machine-learning engineer with published computer-vision research and 9+ years building production software. I take deep-learning and LLM work past the prototype stage: imaging models pulled into hospital workflows, an AI freight dispatcher, and multi-model pipelines running in live SaaS. The full-stack background is what gets the models deployed and keeps them running once they're out.

TECHNICAL SKILLS

AI/ML: PyTorch, TensorFlow, EfficientNet, YOLO, GradCAM++, OpenCV, computer vision, LLM/agent systems, MCP, Claude/GPT/DeepSeek APIs

Languages: Python, C#, C++, TypeScript, JavaScript, Kotlin

ML Tooling & Data: FastAPI, OHIF/DICOM, PostgreSQL, MongoDB, Redis, Prisma, NumPy/pandas

Backend: ASP.NET Core, EF Core, SignalR, Node.js, Bun, NestJS, ElysiaJS

Frontend: React, Next.js, Angular, Blazor, Tailwind CSS

Cloud & DevOps: AWS (S3, Lambda, Cognito, Amplify), Azure, Docker, Kubernetes, GitHub Actions

EDUCATION

Northeastern University, M.S. in Computer Science, Boston, MA

Jan 2024 – May 2026

Suffolk University, B.S. in Computer Science, Boston, MA

Sep 2021 – May 2023

PROFESSIONAL EXPERIENCE

Senior AI/ML Engineer | **EmTech Care Labs** | Portland, ME

Jan 2025 – Jun 2025

- Built the machine-learning features for a HIPAA-compliant Alzheimer's care platform, including a risk model that flags patients showing early signs of decline from daily activity and care-log data so caregivers can step in sooner.
- Added an NLP pipeline that reads free-text clinical notes and pulls out structured details (symptoms, medications, care events), cutting the manual charting load for the care team.
- Stood up the ML infrastructure under HIPAA constraints, with de-identified training data, model serving behind the API, and monitoring. The platform cleared an outside HIPAA readiness review with no critical findings.

Senior Machine Learning Engineer | **AllFactors** | Santa Clara, CA

Jul 2022 – Dec 2023

- Built property-valuation and market-forecasting models on real-estate data that powered the product's analytics, turning raw pricing and inventory feeds into predicted values and trend signals.
- Designed the real-time feature pipeline (SignalR + SQL Server change tracking) that fed live market data to the models and pushed updated predictions to hundreds of concurrent users with sub-second latency.
- Productionized the models into the customer-facing dashboards and added automated tests around the data and prediction flows.

Software Engineer | **Smart Meal Service** | Russia

Sep 2020 – Oct 2021

- Built the software for a self-service food kiosk and robotic cashier in C# / .NET, integrating with payment terminals and hardware over low-level protocols, plus the customer ordering app and an admin dashboard.
- Added automated order routing that sent each order straight to the right kitchen station, and a hardware abstraction layer so one codebase ran across three machine models.

Game Developer | **Pentalight Technology** | Malaysia

Mar 2020 – Feb 2021

- Built the multiplayer networking for a VR smart-city simulation in Unity, syncing every participant's view in real time, plus spatial audio, hand-tracking, and a desktop spectator client.

Earlier: Freelance Developer at Freelancer.com (2019 – 2020), building Python NLP and .NET applications. Game Developer at EC Dev Team, Uzbekistan (2016 – 2019), led development of an RTS mod for Paradox's Clausewitz Engine.

SELECTED AI/ML PROJECTS

MelanomaNet

[GitHub](#) | [Paper](#)

PyTorch, EfficientNet V2, GradCAM++ , OpenCV

- Trained a deep-learning model to classify skin lesions across nine diagnostic categories on a standard public dermatology dataset of 25K+ images, reaching 0.86 weighted F1.
- The research contribution is interpretability: the model produces heatmaps of what it looked at, decomposed along the same ABCDE criteria (Asymmetry, Border, Color, Diameter, Evolution) dermatologists already use, so a clinician can follow and trust the reasoning rather than getting an opaque score. Published on arXiv.

Med Image Scanner

<https://github.com/suxrobgm/med-image-scanner>

FastAPI, PyTorch, OpenCV, OHIF Viewer, DICOM, Next.js

- HIPAA-compliant platform that pulls X-ray, CT, and MRI scans directly from hospital imaging systems and runs PyTorch detection models on them, flagging pneumonia on chest X-rays and intracranial hemorrhage on head CTs.
- Renders model predictions as overlays inside a real radiology viewer (OHIF), and handles patient data safely throughout, with on-the-fly de-identification, audit logging, and role-based access.

Bookshelf Scanner

<https://github.com/suxrobgm/bookshelf-scanner>

YOLO, Moondream2 VLM, FastAPI, Angular

- Two-stage computer-vision pipeline: a YOLO segmentation model isolates each individual book spine from a shelf photo, then a vision-language model reads the title and author off each spine.
- Wrapped it in a web app for uploading photos, correcting misreads, and exporting the recognized library to Goodreads or LibraryThing.

LogisticsX

<https://logisticsx.app> | [GitHub](#)

Claude API, MCP, .NET 10, EF Core, SignalR, Angular 21, Kotlin KMP, PostgreSQL

- Built an AI dispatcher that matches freight loads to trucks, checks federal driver hours-of-service rules, and plans multi-stop routes on its own, turning a manual 15-minute decision into a near-instant one. Runs on the Claude API through a custom tool-use agent with its own set of tools.
- Shipped it as a real product: web portals for dispatchers, customers, and brokers, a cross-platform driver app, and a multi-tenant .NET backend integrated with the major freight load boards and tracking providers.

OGStack

<https://ogstack.dev> | [GitHub](#)

Next.js 16, ElysiaJS, Bun, Satori, Cloudflare R2, Stripe

- Built a multi-model AI layer that routes work across Claude, GPT, and DeepSeek to read a web page and match its branding, then generates an on-brand social-share image, including AI-generated backgrounds.
- Productionized it as a SaaS serving cached images from a global CDN; it generated 10K+ images during the beta.